

Answer 1

I first tried the `.split()` method of python. This gave an accuracy of 44% which is the same as majority prediction. The error was that the `.split()` method only separates at spaces, thereby considering “good” and “good!” as two separate words. This leads to the individual terms to be very low, thereby causing a majority prediction since the sum is not outweighing the difference introduced by $p(y)$.

Then I used `word_tokenize` from `nltk` but that ran into underflow issues (the net size of the emails became very large as it counts multiple punctuations as different words)

Finally I used `CountVectorizer` from `SKLearn`.

Random prediction: 20%

Majority prediction: 44%

My Implementation: 60.035% (test set)
63.05% (training set)

There is a significant improvement over both the random prediction and majority prediction

Confusion matrix:

```
[[ 14335.  2785.  1360.  1080.  3136.]  
 [ 3845.  3290.  1675.   727.   323.]  
 [ 1187.  3335.  5190.  2543.   574.]  
 [  449.  1081.  5410. 18027. 15353.]  
 [  353.   347.   896. 6981. 39436.]]
```

Highest value for diagonal entry: 5

The classification is not very good as the non-diagonal entries are quite high. This is to be expected since the classes themselves have some dependence in reality.

47% accuracy using the function provided (test set) after both stemming and stopword removal

Feature engineering

Bigrams: 63.28 % on test set

```
[[ 1.66810000e+04  4.11900000e+03  1.76700000e+03  8.37000000e+02  
   1.42700000e+03]  
 [ 9.12000000e+02  7.90000000e+02  2.14000000e+02  2.10000000e+01  
   2.00000000e+00]  
 [ 7.57000000e+02  1.84100000e+03  1.69300000e+03  2.40000000e+02  
   4.00000000e+01]]
```

```
[ 1.30900000e+03  3.49800000e+03  9.53100000e+03  1.86250000e+04
 1.05240000e+04]
[ 5.10000000e+02  5.90000000e+02  1.32600000e+03  9.63500000e+03
 4.68290000e+04]]
```

POS+Bigrams=63.19%

```
[[ 1.68110000e+04  4.22800000e+03  1.75000000e+03  7.79000000e+02
 1.23700000e+03]
[ 5.80000000e+02  5.05000000e+02  1.32000000e+02  9.00000000e+00
 1.00000000e+00]
[ 5.74000000e+02  1.35800000e+03  1.17100000e+03  1.58000000e+02
 2.90000000e+01]
[ 1.55800000e+03  3.99200000e+03  9.87600000e+03  1.77070000e+04
 9.24700000e+03]
[ 6.46000000e+02  7.55000000e+02  1.60200000e+03  1.07050000e+04
 4.83080000e+04]]
```

F1 score:

1:0.65
2:0.34
3:0.36
4:0.54
5:0.74

Macro avg:0.6

Micro avg:0.52

Weighted avg:0.61

Good metric to see the problem with majority prediction

Full training set:

accuracy:72.021%

F1 value:0.67